



# ScalaMeter

Performance regression testing framework





Aleksandar Prokopec, Josh Suereth, Vlad Ureche, Roman Zoller, Ngoc  
Duy Pham, Alexey Romanov, Roger Vion, Eugene Platonov, Lukas Rytz,  
Paolo Giarusso, Dan Burkert, Christian Krause, and others

# Background



# Background



Georges A., Buytaert D., Eeckhout L.  
Statistically rigorous Java performance evaluation. **OOPSLA '07**

# Motivation

```
List(1 to 100000: _*).map(x => x * x)
```

26 ms

# Motivation

```
List(1 to 100000: _*).map(x => x * x)
```

26 ms

```
class List[+T]  
  extends Seq[T] {  
  
    // implementation 1  
  
}
```

# Motivation

```
List(1 to 100000: _*).map(x => x * x)
```

26 ms



```
class List[+T]  
extends Seq[T] {
```

```
// implementation 1
```

```
}
```

# Motivation

```
List(1 to 100000: _*).map(x => x * x)
```

49 ms



```
class List[+T]  
extends Seq[T] {
```

```
// implementation 2
```

```
}
```

# First example

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```



# First example

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

# First example

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

# First example

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

# First example

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}  
measure()
```

# First example

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}  
measure()
```

26 ms

# The warmup problem

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}  
measure()  
measure()
```

26 ms, 11 ms

# The warmup problem

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}  
measure()  
measure()
```

Why?

Mainly:

- JIT compilation
- dynamic optimization

26 ms, 11 ms

# The warmup problem

```
def measure() {  
  val buffer = mutable.ArrayBuffer(0 until 2000000: _*)  
  val start = System.currentTimeMillis()  
  var sum = 0  
  buffer.foreach(sum += _)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}  
measure()  
measure()
```

26 ms, 11 ms

45 ms, 10 ms



# The warmup problem

```
def measure2() {  
  val buffer = mutable.ArrayBuffer(0 until 4000000: _*)  
  val start = System.currentTimeMillis()  
  buffer.map(_ + 1)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

# The warmup problem

```
def measure2() {  
  val buffer = mutable.ArrayBuffer(0 until 4000000: _*)  
  val start = System.currentTimeMillis()  
  buffer.map(_ + 1)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

241, 238, 235, 236, 234

# The warmup problem

```
def measure2() {  
  val buffer = mutable.ArrayBuffer(0 until 4000000: _*)  
  val start = System.currentTimeMillis()  
  buffer.map(_ + 1)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

241, 238, 235, 236, 234, 429

# The warmup problem

```
def measure2() {  
  val buffer = mutable.ArrayBuffer(0 until 4000000: _*)  
  val start = System.currentTimeMillis()  
  buffer.map(_ + 1)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

241, 238, 235, 236, 234, 429, 209

# The warmup problem

```
def measure2() {  
  val buffer = mutable.ArrayBuffer(0 until 4000000: _*)  
  val start = System.currentTimeMillis()  
  buffer.map(_ + 1)  
  val end = System.currentTimeMillis()  
  println(end - start)  
}
```

241, 238, 235, 236, 234, 429, 209, 194, 195, 195


# The warmup problem

Bottomline: benchmark has to be repeated until the running time becomes “stable”.

The number of repetitions is not known in advance.


241, 238, 235, 236, 234, 429, 209, 194, 195, 195

# Warming up the JVM



241, 238, 235, 236, 234, 429, 209, 194, 195, 195,  
194, 194, 193, 194, 196, 195, 195

# Warming up the JVM



241, 238, 235, 236, 234, 429, 209, 194, 195, 195,  
194, 194, 193, 194, 196, 195, 195



# The interference problem

```
val buffer = ArrayBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

```
val buffer = ListBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

# The interference problem

```
val buffer = ArrayBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

```
val buffer = ListBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

Lets measure the first **map** 3 times with 7 repetitions:

61, 54, 54, 54, 55, 55, 56

186, 54, 54, 54, 55, 54, 53

54, 54, 53, 53, 53, 54, 51

# The interference problem

```
val buffer = ArrayBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

```
val buffer = ListBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

Now, lets measure the **list buffer map** in between:

61, 54, 54, 54, 55, 55, 56

186, 54, 54, 54, 55, 54, 53

54, 54, 53, 53, 53, 54, 51

59, 54, 54, 54, 54, 54, 54

44, 36, 36, 36, 35, 36, 36

45, 45, 45, 45, 44, 46, 45

18, 17, 18, 18, 17, 292, 16

45, 45, 44, 44, 45, 45, 44

# The interference problem

```
val buffer = ArrayBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

```
val buffer = ListBuffer(0 until 900000: _*)  
buffer.map(_ + 1)
```

Now, lets measure the **list buffer map** in between:

61, 54, 54, 54, 55, 55, 56

186, 54, 54, 54, 55, 54, 53

54, 54, 53, 53, 53, 54, 51

59, 54, 54, 54, 54, 54, 54

44, 36, 36, 36, 35, 36, 36

45, 45, 45, 45, 44, 46, 45

18, 17, 18, 18, 17, 292, 16

45, 45, 44, 44, 45, 45, 44

# Using separate JVM

Bottomline: always run the tests in a new JVM.

# Using separate JVM

Bottomline: always run the tests in a new JVM.

This may not reflect a real-world scenario, but it gives a good idea of how different several alternatives are.

# Using separate JVM

Bottomline: always run the tests in a new JVM.

It results in a reproducible, more stable measurement.

# The List.map example

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```



# The List.map example

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

37, 38, 37, 1175, 38, 37, 37, 37, 37, ..., 38, 37, 37,  
37, 37, 465, 35, 35, ...

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

This benchmark triggers GC cycles!

37, 38, 37, 1175, 38, 37, 37, 37, 37, ..., 38, 37, 37,  
37, 37, 465, 35, 35, ...

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

This benchmark triggers GC cycles!

37, 38, 37, 1175, 38, 37, 37, 37, 37, ..., 38, 37, 37,  
37, 37, 465, 35, 35, ... -> mean: 47 ms

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

This benchmark triggers GC cycles!

37, 38, 37, 1175, 38, 37, 37, 37, 37, ..., 38, 37, 37,  
37, 37, 465, 35, 35, ... -> mean: 47 ms

37, 37, 37, 647, 37, 36, 38, 37, 36, ..., 36, 37, 36,  
37, 36, 37, 534, 36, 33, ... -> mean: 39 ms

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

Solutions:

- repeat A LOT of times –an accurate mean, but takes A LONG time

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

## Solutions:

- repeat A LOT of times –an accurate mean, but takes A LONG time
- ignore the measurements with GC – gives a reproducible value, and less measurements

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

## Solutions:

- repeat A LOT of times –an accurate mean, but takes A LONG time
- ignore the measurements with GC – gives a reproducible value, and less measurements
  - how to do this?

# The garbage collection problem

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

- manually - verbose:gc



# Automatic GC detection

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

- manually - `verbose:gc`
- automatically using callbacks in JDK7

37, 37, 37, 647, 37, 36, 38, 37, 36, ..., 36, 37, 36,  
37, 36, 37, 534, 36, 33, ...

# Automatic GC detection

```
val list = (0 until 2500000).toList  
list.map(_ % 2 == 0)
```

- manually - `verbose:gc`
- automatically using callbacks in JDK7

raises a GC event



37, 37, 37, 647, 37, 36, 38, 37, 36, ..., 36, 37, 36,  
37, 36, 37, 534, 36, 33, ...

# The runtime problem

- there are other runtime events beside GC – e.g. JIT compilation, dynamic optimization, etc.
- these take time, but cannot be determined accurately

# The runtime problem

- there are other runtime events beside GC – e.g. JIT compilation, dynamic optimization, etc.
- these take time, but cannot be determined accurately
- heap state also influences memory allocation patterns and performance

# The runtime problem

- there are other runtime events beside GC – e.g. JIT compilation, dynamic optimization, etc.
- these take time, but cannot be determined accurately
- heap state also influences memory allocation patterns and performance

```
val list = (0 until 4000000).toList  
list.groupBy(_ % 10)
```

(allocation intensive)

# The runtime problem

- there are other runtime events beside GC – e.g. JIT compilation, dynamic optimization, etc.
- these take time, but cannot be determined accurately
- heap state also influences memory allocation patterns and performance

```
val list = (0 until 4000000).toList  
list.groupBy(_ % 10)
```

120, 121, 122, 118, 123, 794, 109, 111, 115, 113, 110

# The runtime problem

- there are other runtime events beside GC – e.g. JIT compilation, dynamic optimization, etc.
- these take time, but cannot be determined accurately
- heap state also influences memory allocation patterns and performance

```
val list = (0 until 4000000).toList  
list.groupBy(_ % 10)
```

affects the mean – 116 ms vs 178 ms

120, 121, 122, 118, 123,  794, 109, 111, 115, 113, 110

# Outlier elimination

120, 121, 122, 118, 123, 794, 109, 111, 115, 113, 110



# Outlier elimination

120, 121, 122, 118, 123, 794, 109, 111, 115, 113, 110



sort

109, 110, 111, 113, 115, 118, 120, 121, 122, 123, 794

# Outlier elimination

120, 121, 122, 118, 123, 794, 109, 111, 115, 113, 110



sort

109, 110, 111, 113, 115, 118, 120, 121, 122, 123, 794



inspect tail and its variance contribution

109, 110, 111, 113, 115, 118, 120, 121, 122, 123

# Outlier elimination

120, 121, 122, 118, 123, 794, 109, 111, 115, 113, 110



sort

109, 110, 111, 113, 115, 118, 120, 121, 122, 123, 794



inspect tail and its variance contribution

109, 110, 111, 113, 115, 118, 120, 121, 122, 123



redo the measurement

109, 110, 111, 113, 115, 118, 120, 121, 122, 123, 124

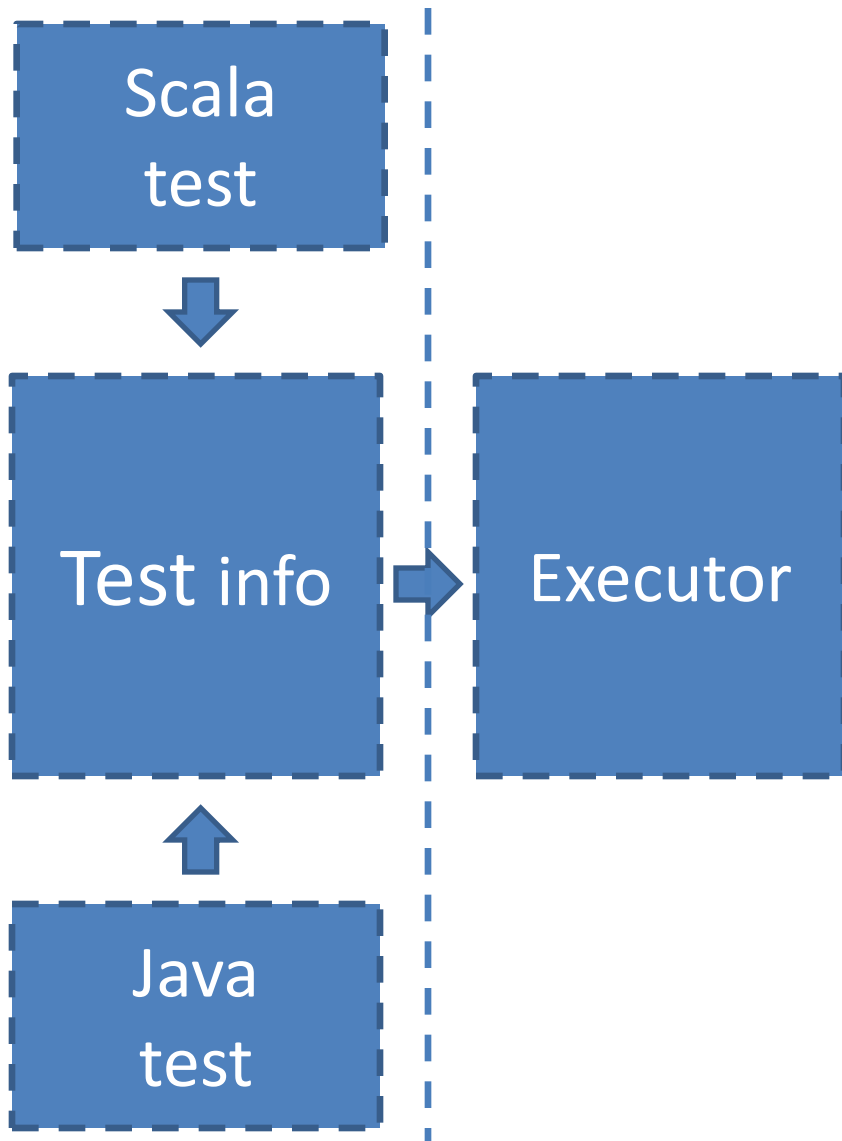
# ScalaMeter



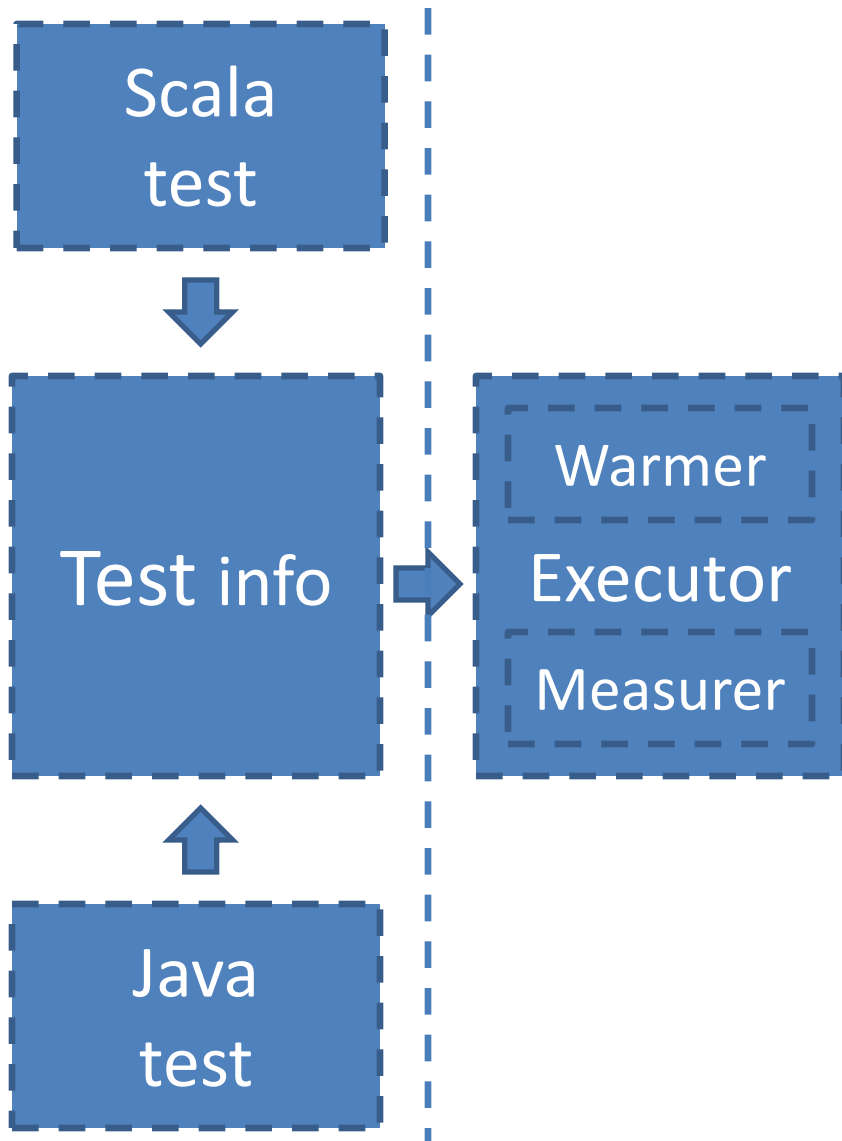
# ScalaMeter



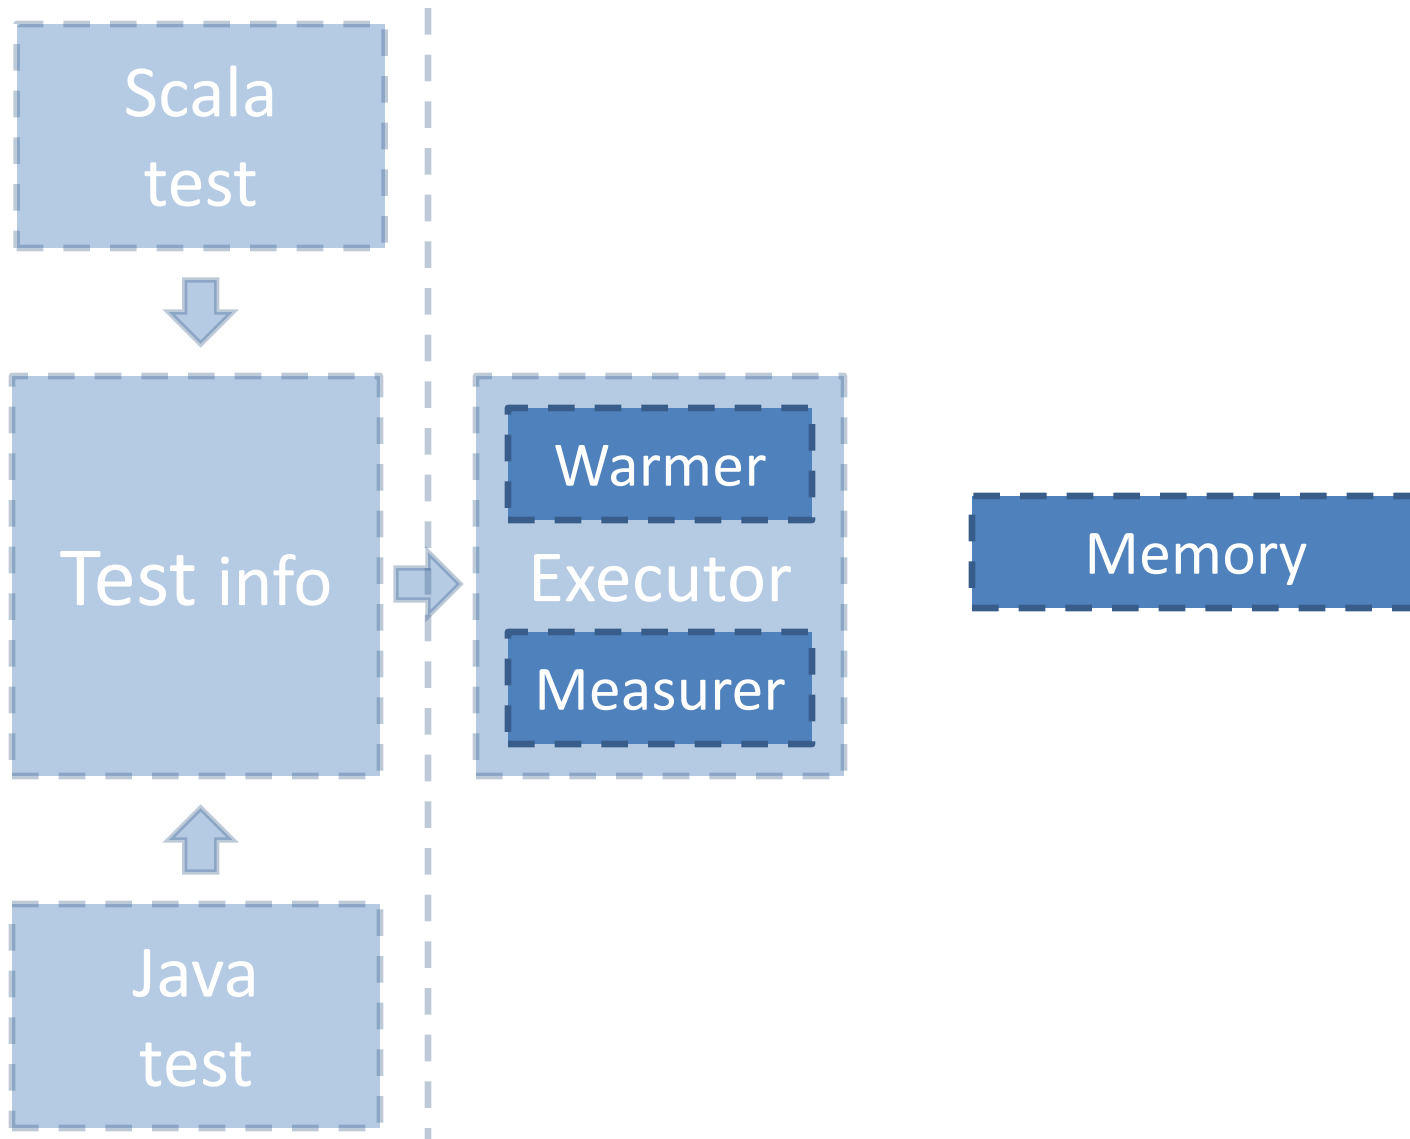
# ScalaMeter



# ScalaMeter

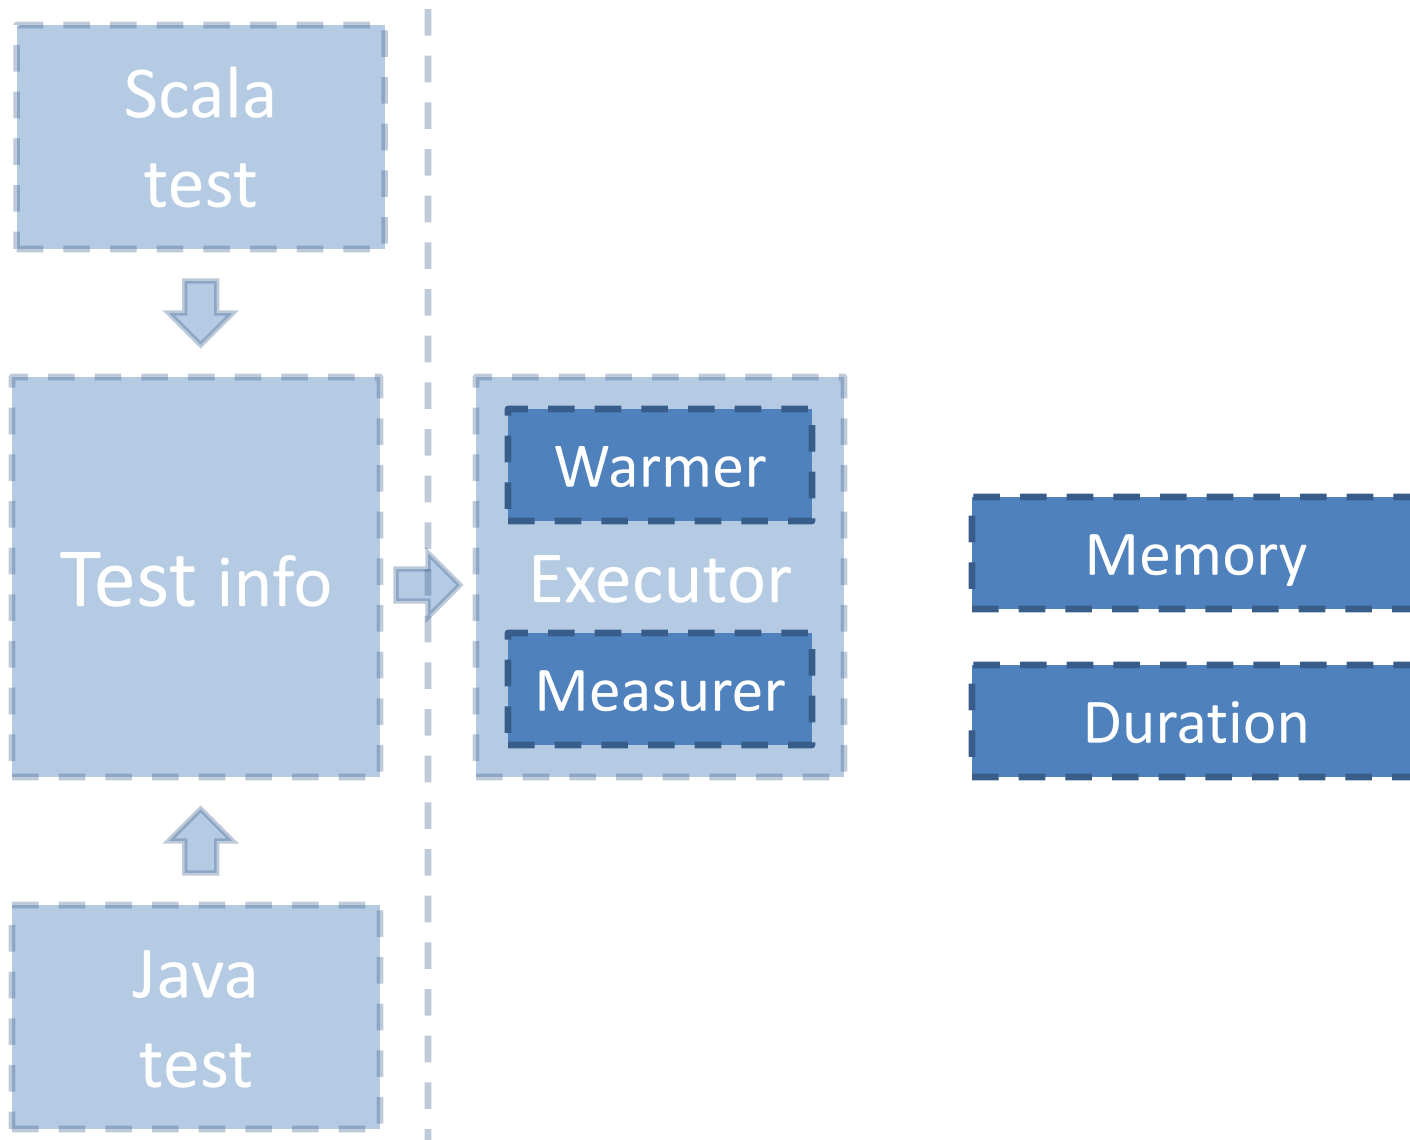


# ScalaMeter

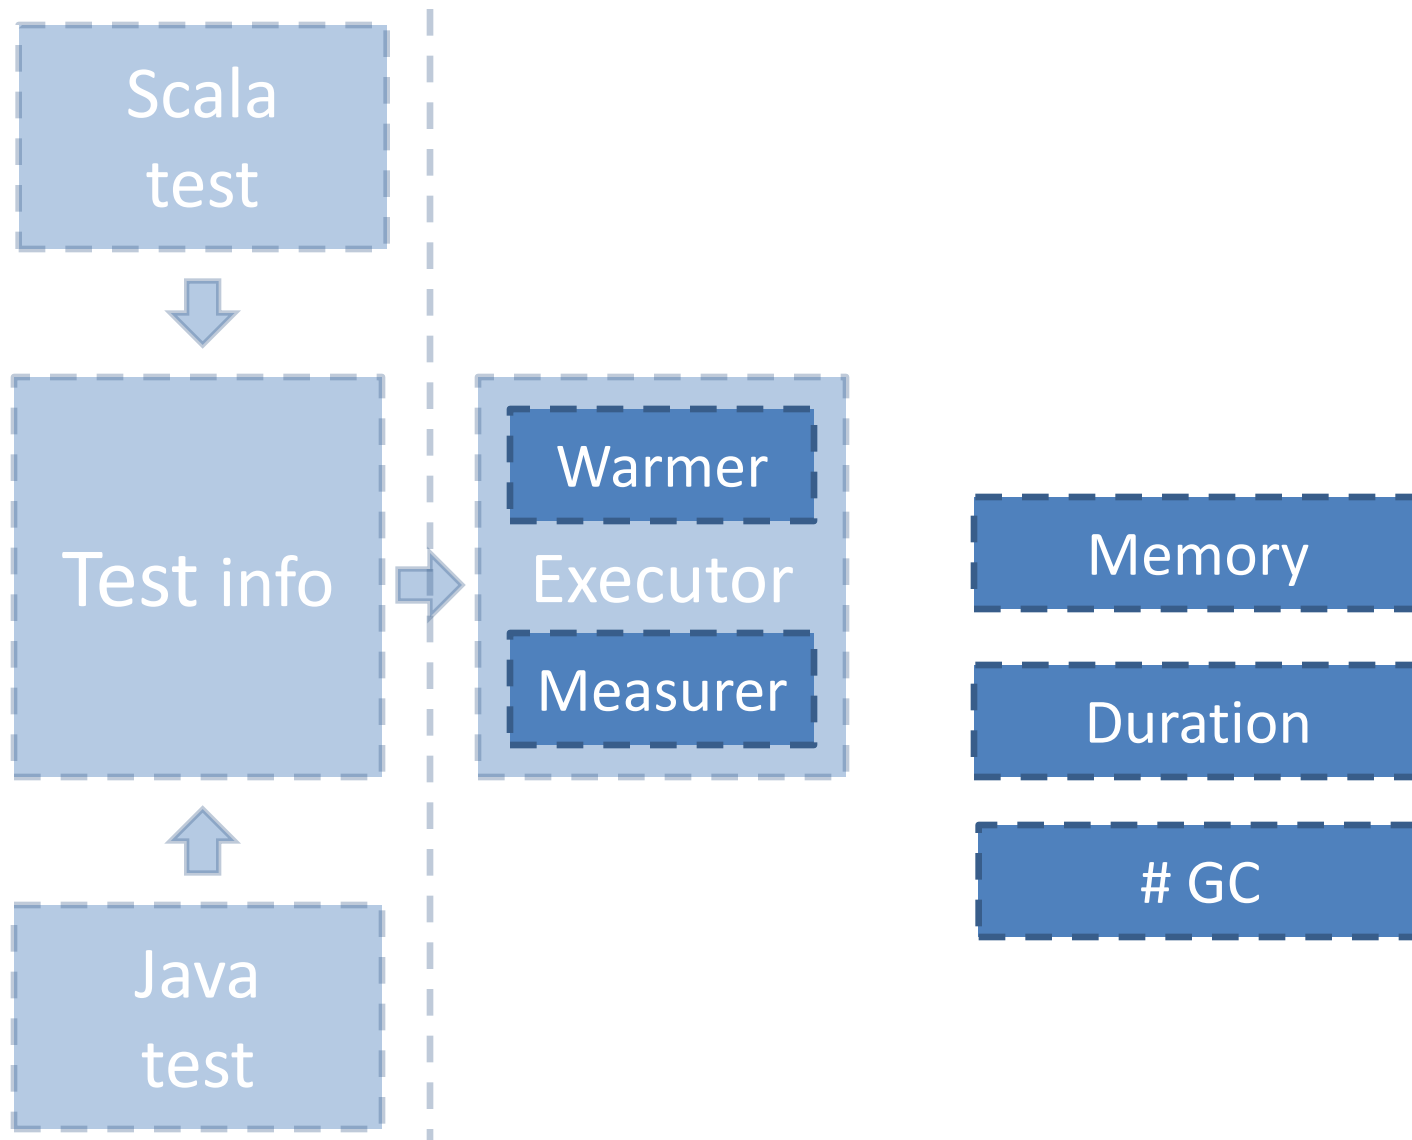




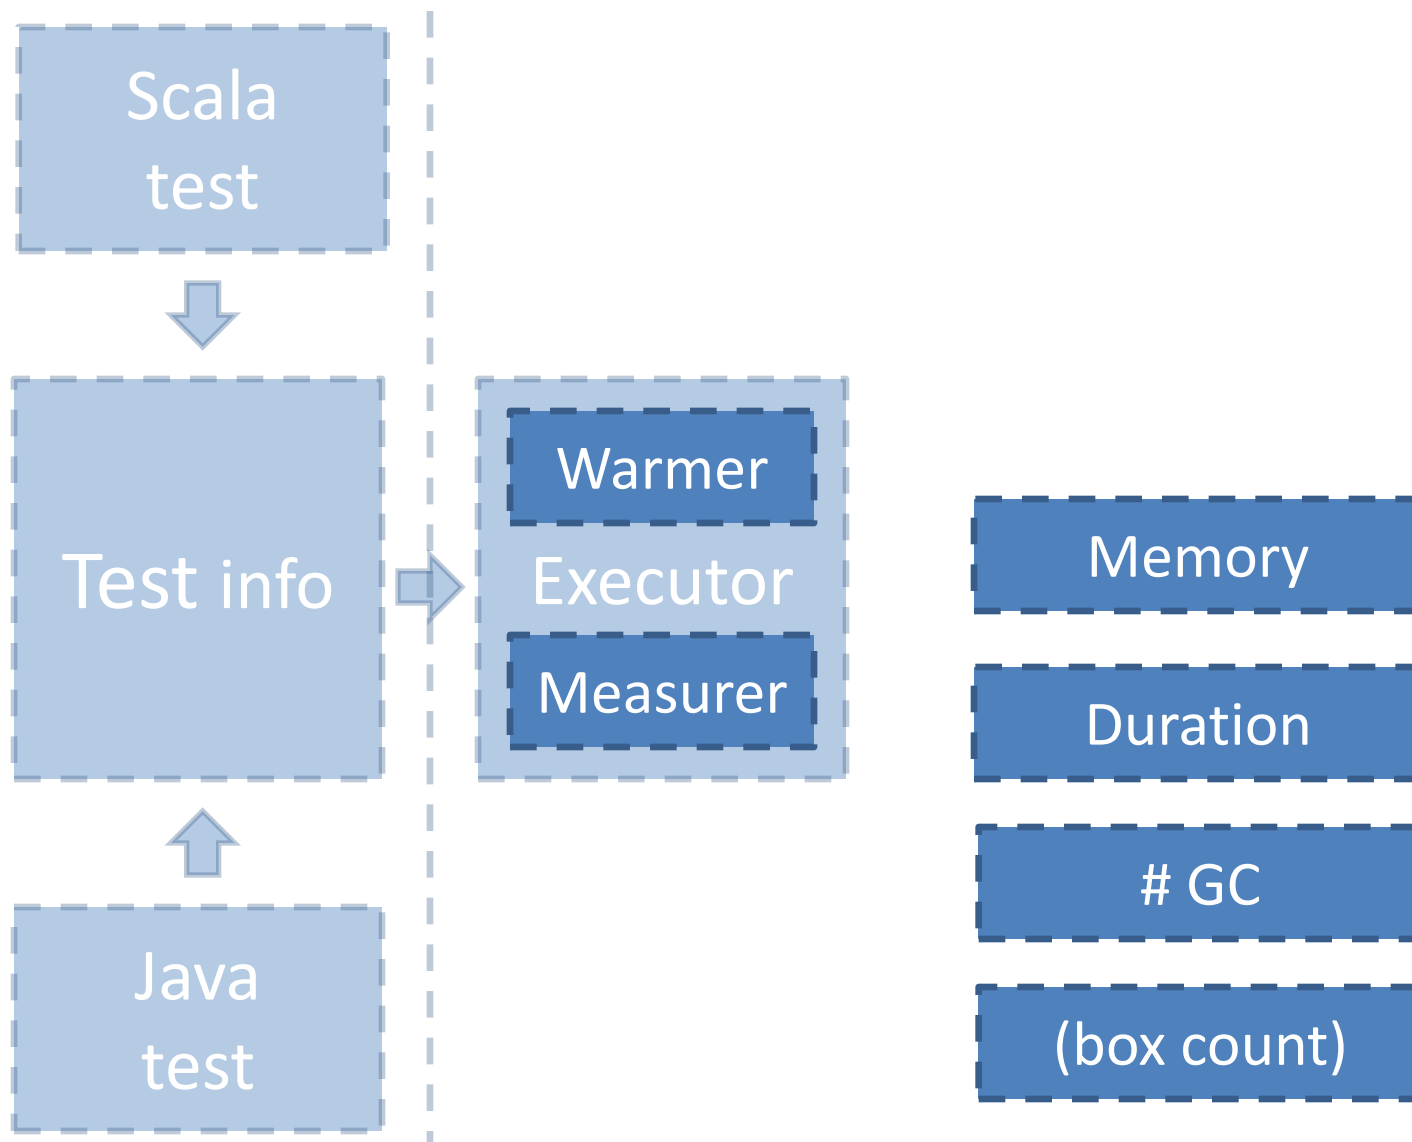
# ScalaMeter



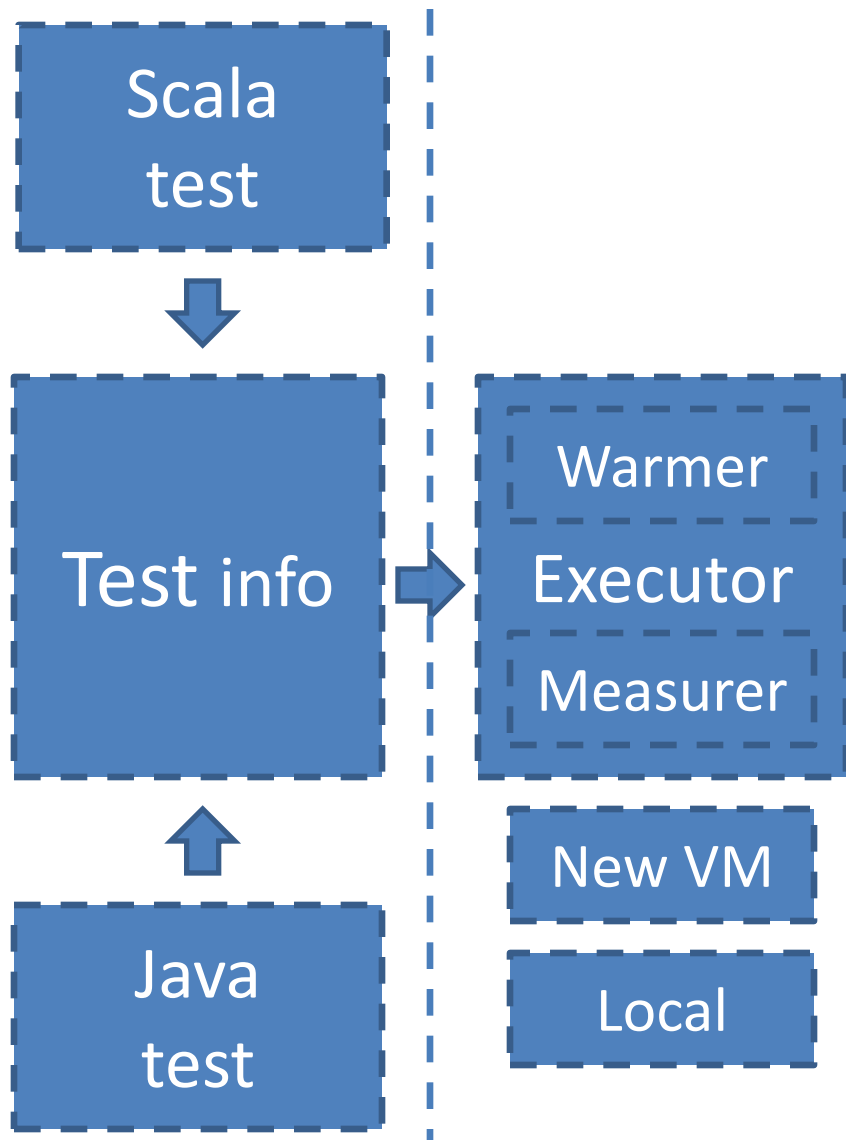
# ScalaMeter



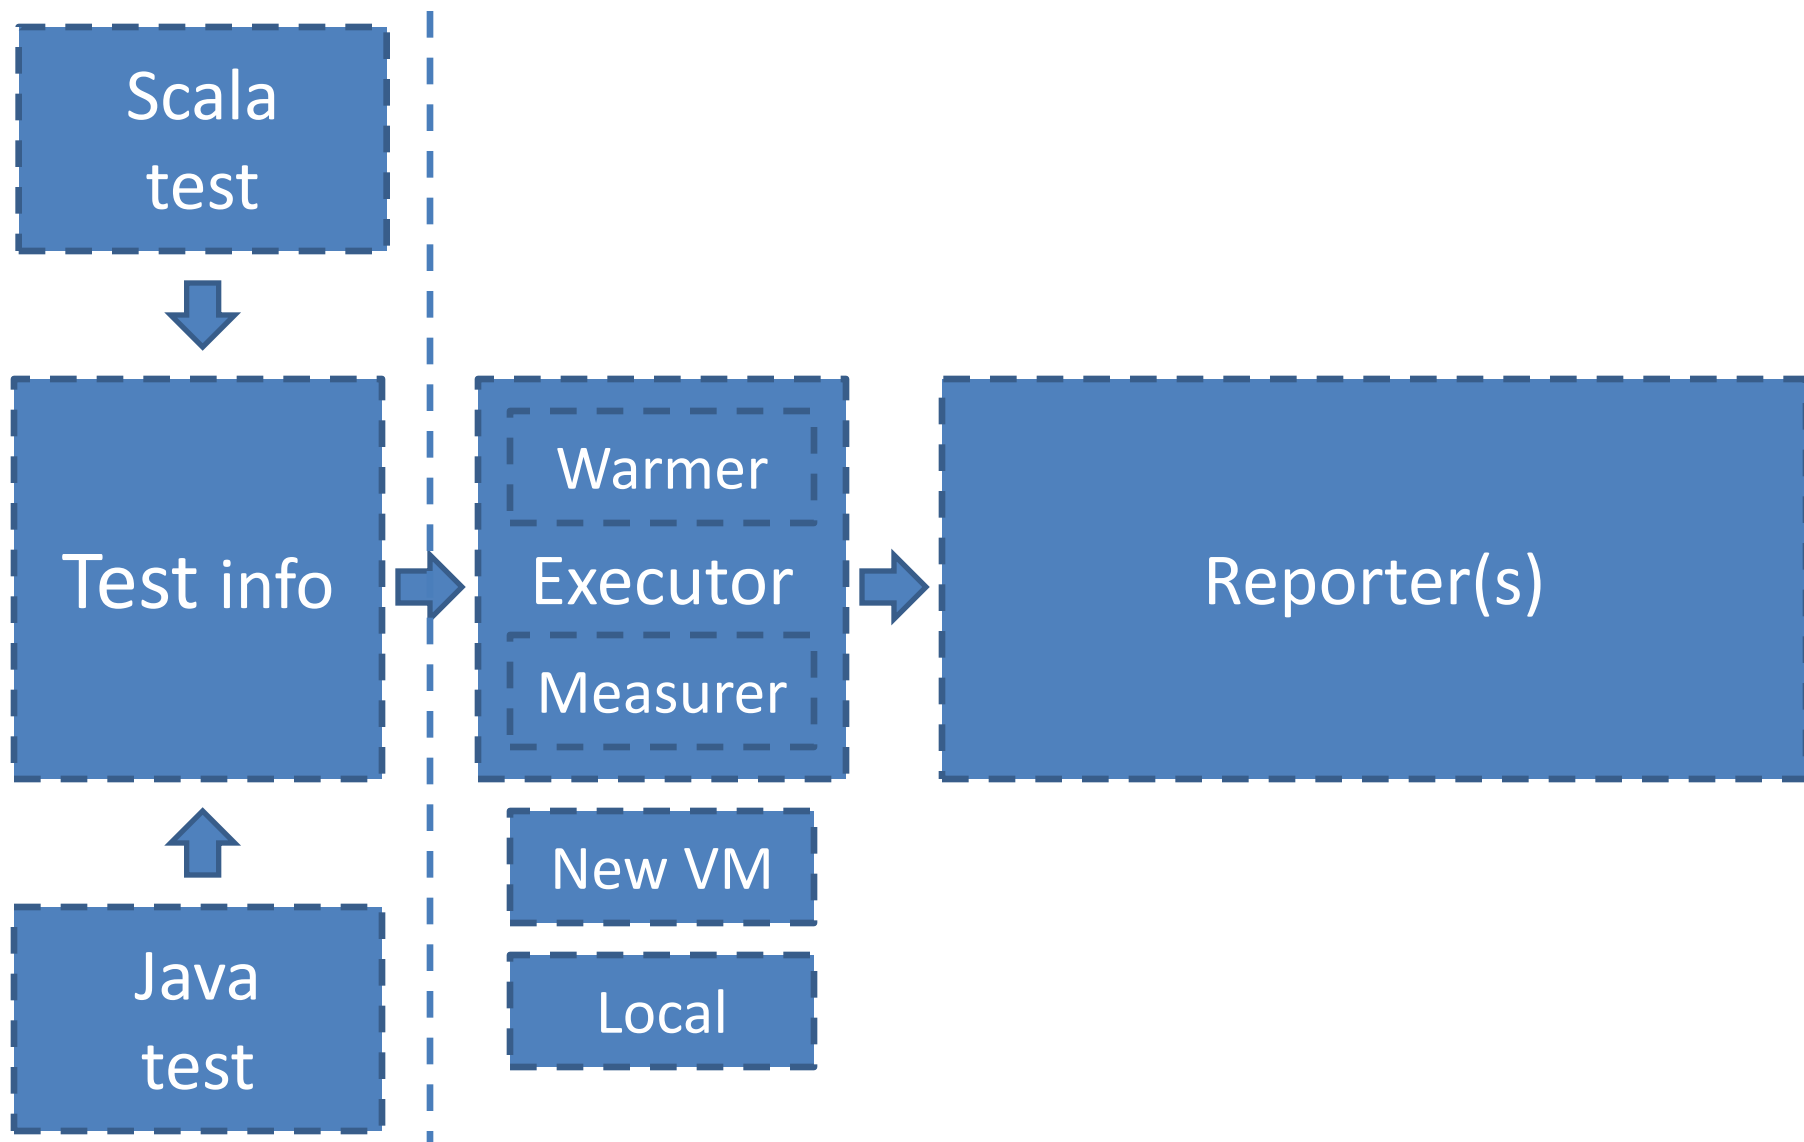
# ScalaMeter



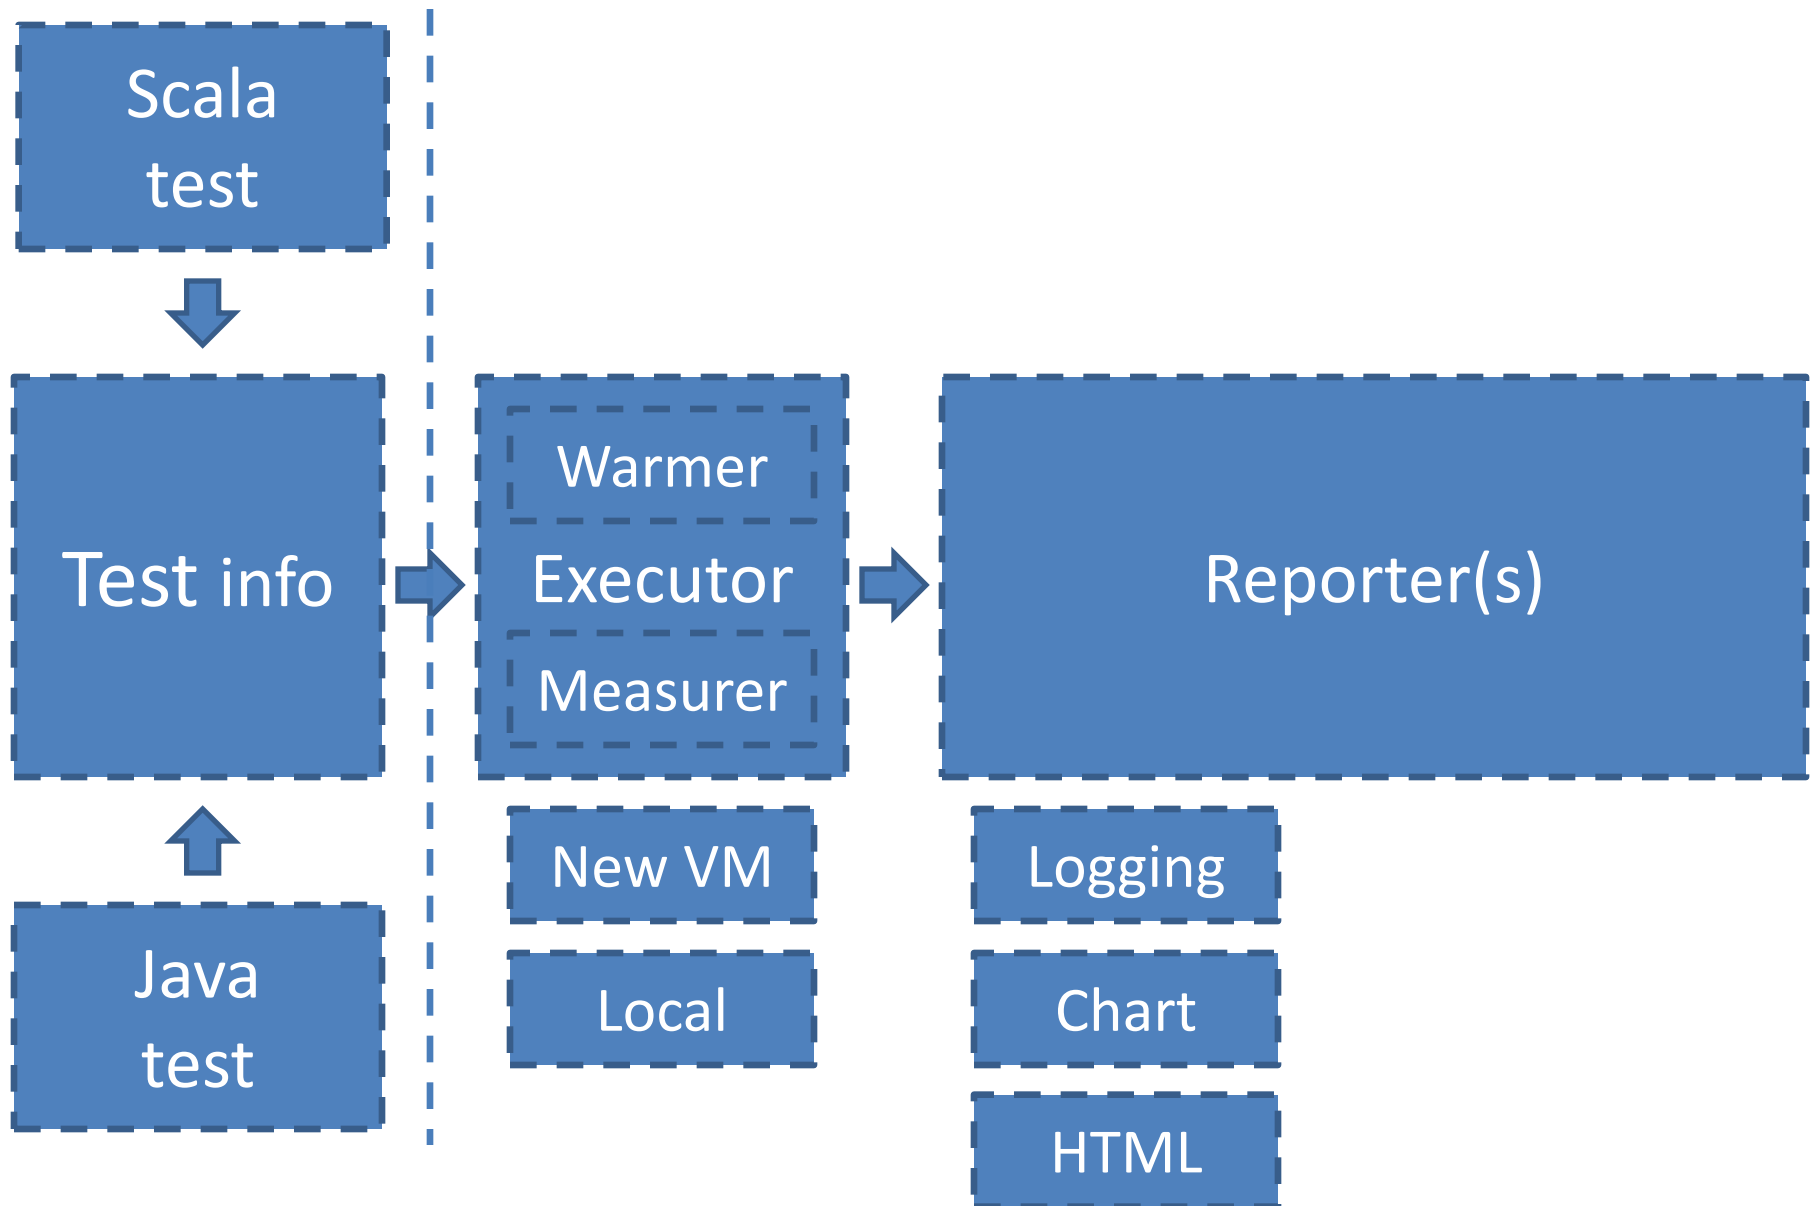
# ScalaMeter



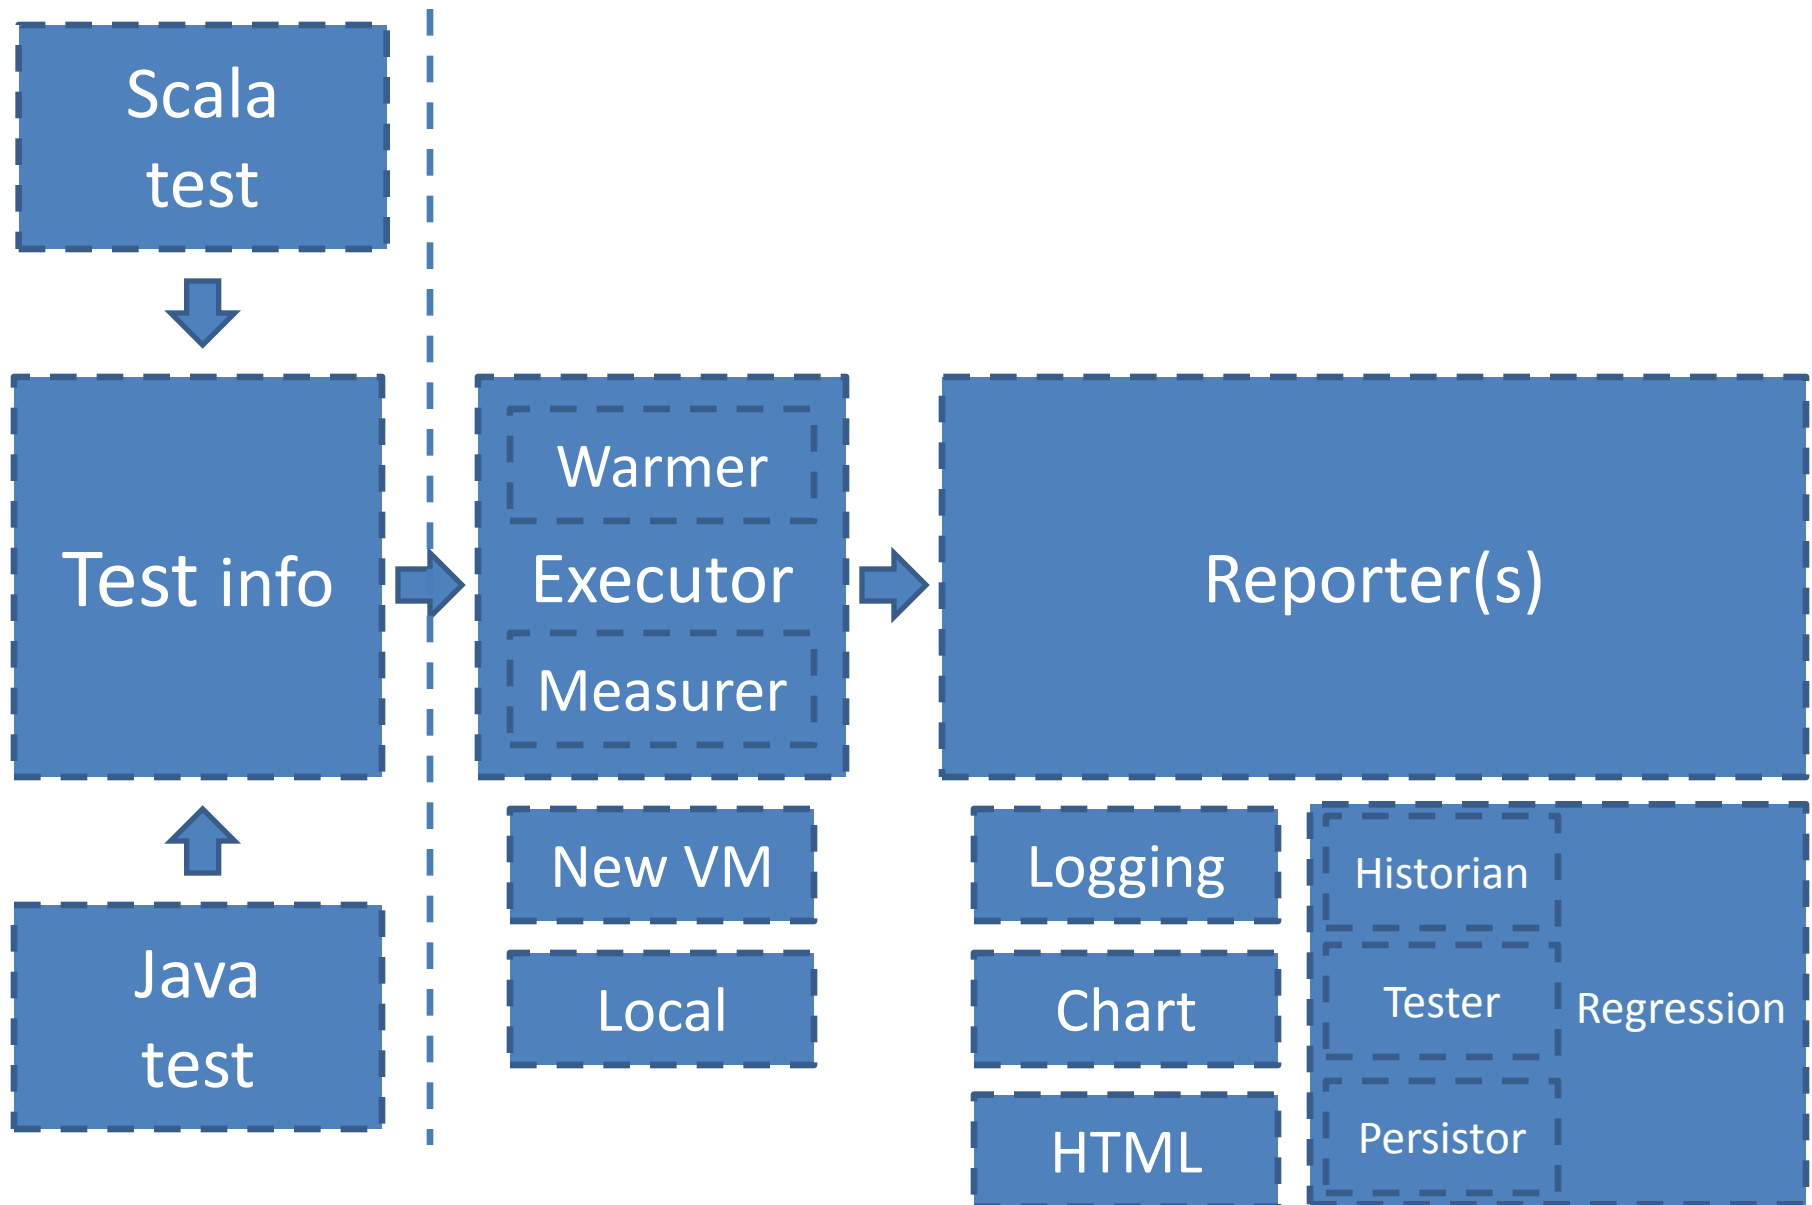
# ScalaMeter



# ScalaMeter



# ScalaMeter



# ScalaMeter example

Scala  
test

```
object ListTest extends PerformanceTest.Microbenchmark {
```

A range of predefined benchmark types



# ScalaMeter example

Scala  
test

Generator(s)

```
object ListTest extends PerformanceTest.Microbenchmark {  
  val sizes = Gen.range("size")(500000, 1000000, 100000)
```

Generators provide input data for tests

# ScalaMeter example

Scala  
test

Generator(s)

```
object ListTest extends PerformanceTest.Microbenchmark {  
  val sizes = Gen.range("size")(500000, 1000000, 100000)  
  val lists = for (sz <- sizes) yield (0 until sz).toList
```

Generators can be composed a la ScalaCheck

# ScalaMeter example

Scala  
test

Generator(s)

Snippet(s)

```
object ListTest extends PerformanceTest.Microbenchmark {  
  val sizes = Gen.range("size")(500000, 1000000, 100000)  
  val lists = for (sz <- sizes) yield (0 until sz).toList  
  
  using(lists) in { xs =>  
    xs.groupBy(_ % 10)  
  }  
}
```

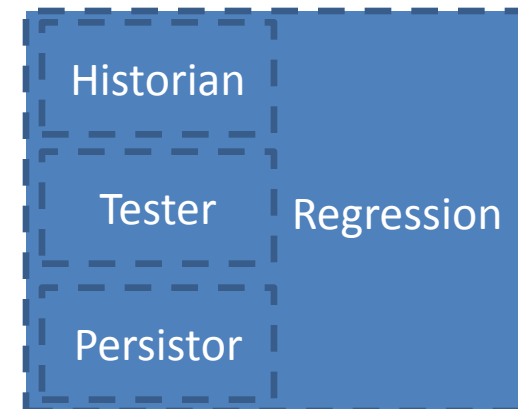
Concise syntax to specify and group tests

# ScalaMeter example

```
object ListTest extends PerformanceTest.Microbenchmark {  
  val sizes = Gen.range("size")(500000, 1000000, 100000)  
  val lists = for (sz <- sizes) yield (0 until sz).toList  
  
  measure method "groupBy" in {  
    using(lists) in { xs =>  
      xs.groupBy(_ % 10)  
    }  
  
    using(ranges) in { xs =>  
      xs.groupBy(_ % 10)  
    }  
  }  
}
```

# Automatic regression testing

```
using(lists) in { xs =>  
  var sum = 0  
  xs.foreach(x => sum += x)  
}
```



# Automatic regression testing


```
using(lists) in { xs =>  
  var sum = 0  
  xs.foreach(x => sum += x)  
}
```

[info] Test group: foreach

[info] - foreach.Test-0 measurements:


[info] - at size -> 2000000, 1 alternatives: passed

[info] (ci = <7.28, 8.22>, significance = 1.0E-10)



# Automatic regression testing

```
using(lists) in { xs =>  
  var sum = 0  
  xs.foreach(x => sum += math.sqrt(x))  
}
```



# Automatic regression testing

```
using(lists) in { xs =>  
  var sum = 0  
  xs.foreach(x => sum += math.sqrt(x))  
}
```

[info] Test group: foreach

[info] - foreach.Test-0 measurements:

[info] - at size -> 2000000, 2 alternatives: failed

[info] (ci = <14.57, 15.38>, significance = 1.0E-10)

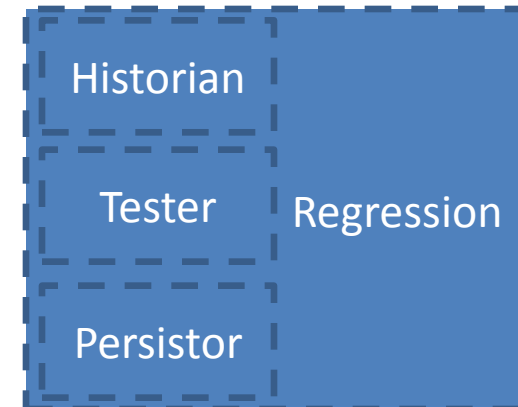

[error] Failed confidence interval test: <-7.85, -6.60>

[error] Previous (mean = 7.75, stdev = 0.44, ci = <7.28, 8.22>)

[error] Latest (mean = 14.97, stdev = 0.38, ci = <14.57, 15.38>)



# Automatic regression testing



# Report generation

<http://scala-blitz.github.io/home/documentation/benchmarks//chara.html>

# Online mode

```
import org.scalameter._  
val time = measure {  
  for (i <- 0 until 100000) yield i  
}  
println(s"Total time: $time")
```

# Online mode

```
val time = config(
  Key.exec.benchRuns -> 20,
  Key.verbose -> true
) withWarmer {
  new Warmer.Default
} withMeasurer {
  new Measurer.IgnoringGC
} measure {
  for (i <- 0 until 100000) yield i
}
println(s"Total time: $time")
```

# Tutorials online!

<http://scalameter.github.io>

Thank you!