Advanced Computing Laboratory: Three Research Threads

Markus Püschel
Computer Science

ETH zürich
SPIRAL
www.spiral.net


Marcela Zuluaga


Georg Ofenbeck


Victoria Caparròs Cabezas

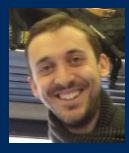
Daniele Spampinato


Alen Stojanov


François Serre


Gagandeep Singh

1. **Program synthesis for performance**
   Daniele Spampinato [CGO 2014]
2. **Performance bottleneck modeling**
   Victoria Caparròs Cabezas [IISWC 2014]
3. **Predicting Pareto-optimal solutions**
   Marcela Zuluaga, Andreas Krause [ICML 2013]

Vision:
Program Synthesis For Performance

Generate highest performance code for mathematical computations directly from a mathematical description

Approach

Mathematical DSLs
Rewriting systems for difficult optimizations
Compiler
Learning and search for fine-tuning

Use advanced software platforms for the development of generators

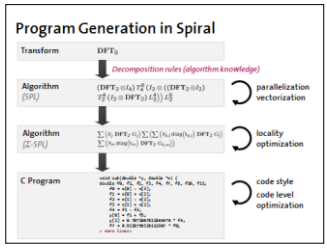

Example: Linear transforms
www.spiral.net

LGen: Generator for Linear Algebra

$$\gamma = x^T(A+B)y + \delta$$ ← A is 2x3, **x** is 3x1, ...

LGen *Design similar to Spiral*

```
void f(double const * A, double const * x, double * y) {
  __m128d t0, …;

  t0 = _mm_loadu_pd(A);
  t1 = _mm_load_sd(A + 2);
  ...
  t6 = _mm_hadd_pd(_mm_mul_pd(t0, t4), _mm_mul_pd(t2, t4));
  t7 = _mm_shuffle_pd(t1, t3, 0);
  t8 = _mm_mul_pd(t7, _mm_shuffle_pd(t5, t5, 0));
  t9 = _mm_add_pd(t6, t8);

  _mm_storeu_pd(y, t9);
}
```

Architecture of LGen

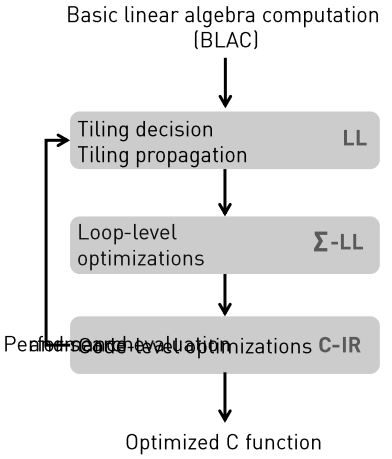

$$y = Ax$$

$$[y = Ax]_{2,1}$$

$$y = \sum_{i,j} S_i(G_i \cdots)$$

```
...
Mov (mmMulPs A[0,0], x[0,0]), t[0,0]
...

for(int i = … ) {
  …
  t = _mm_mul_ps(a, x);
  …
}
```

Example Benchmarks

Intel Xeon Westmere

$$C = \alpha AB + \beta C$$

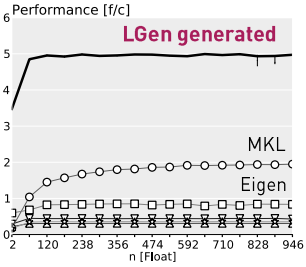

$$A \in \mathbb{R}^{n \times 4}$$

$$B \in \mathbb{R}^{4 \times 4}$$

Intel Xeon Westmere

$$C = \alpha(A_0 + A_1)^T B + \beta C$$

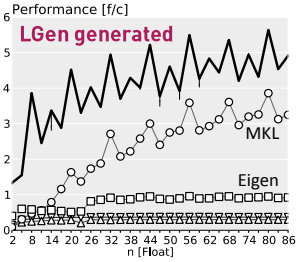

$$A_0 \in \mathbb{R}^{4 \times n}$$

$$B \in \mathbb{R}^{4 \times n}$$

ARM Cortex-A8

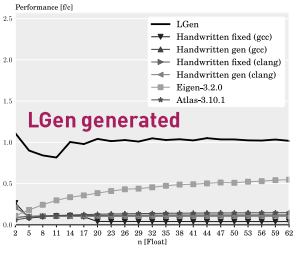

Next Steps

Structured matrices

Higher level algorithms (matrix factorizations etc.)

Multicore

Domain specific extensions

- *Optimization*
- *Machine learning*
- *Communication & Control*

1. **Program synthesis for performance**
   Daniele Spampinato [CGO 2014]
2. **Performance bottleneck modeling**
   Victoria Caparròs Cabezas [IISWC 2014]
3. **Predicting Pareto-optimal solutions**
   Marcela Zuluaga, Andreas Krause [ICML 2013]

Measuring Runtime

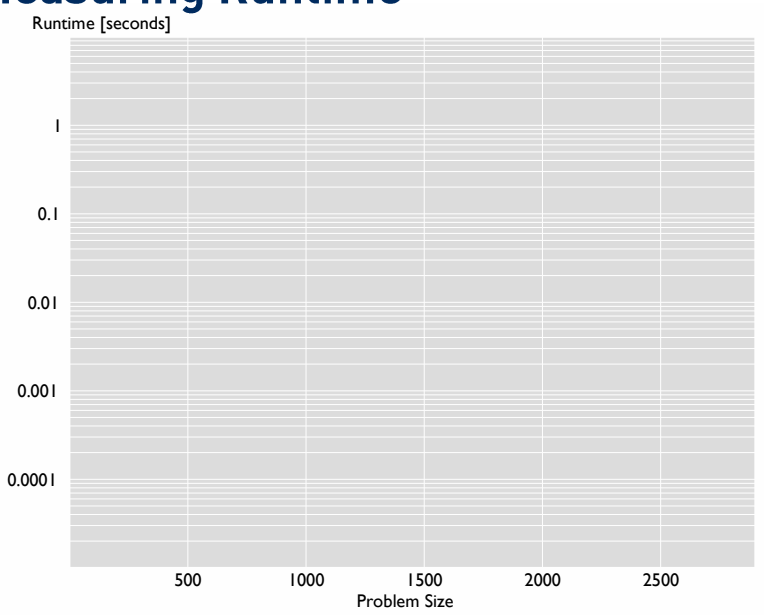

Measuring Runtime

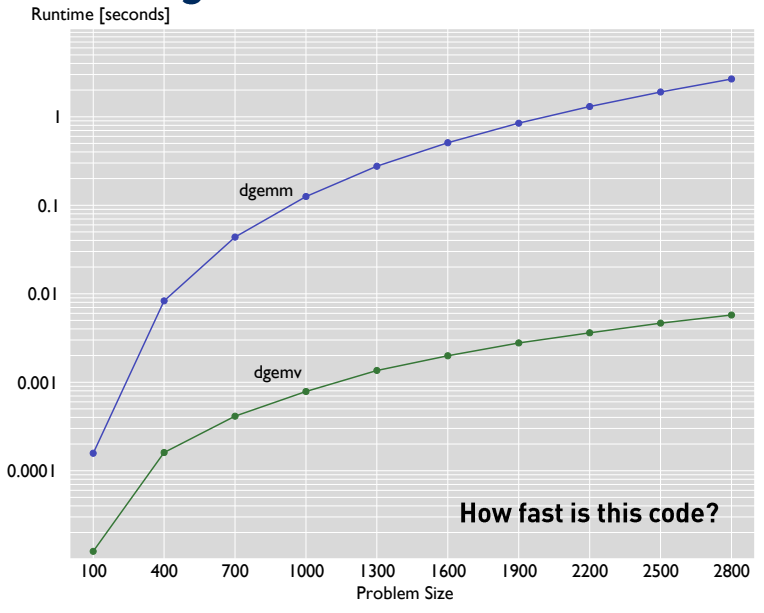

Measuring Performance

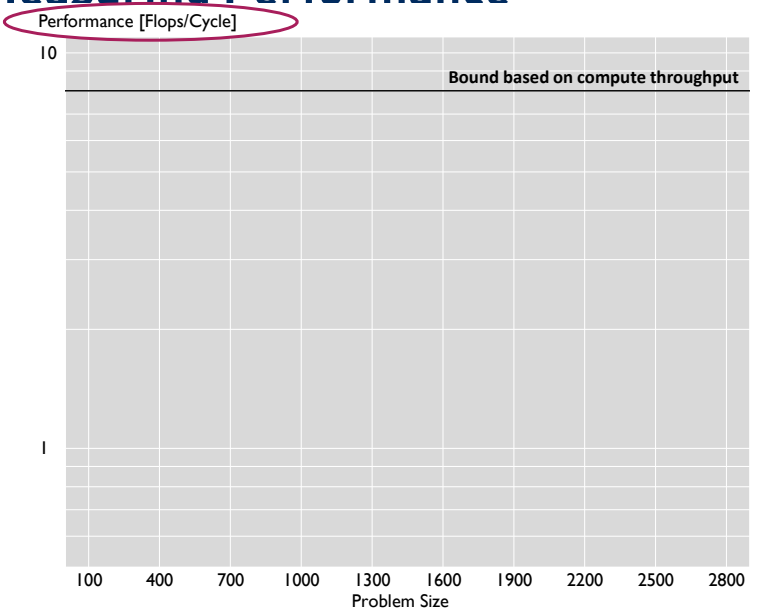

Measuring Performance

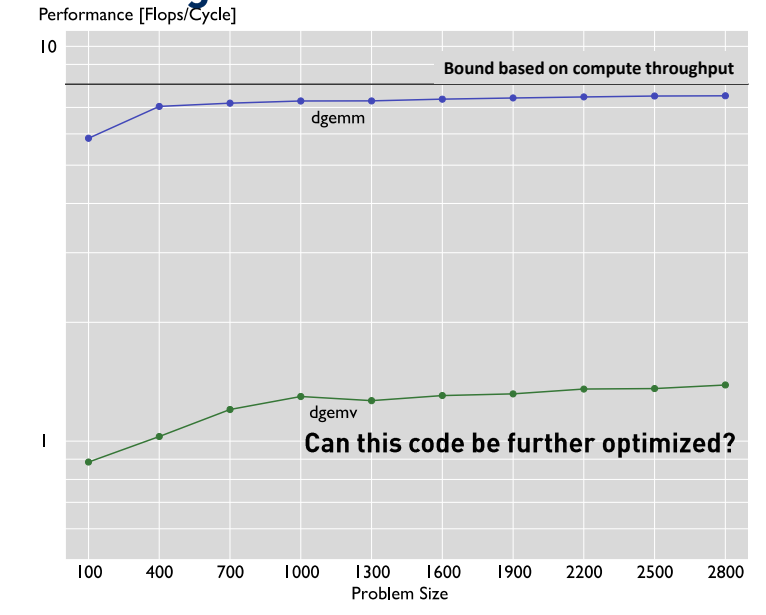

Roofline Plot

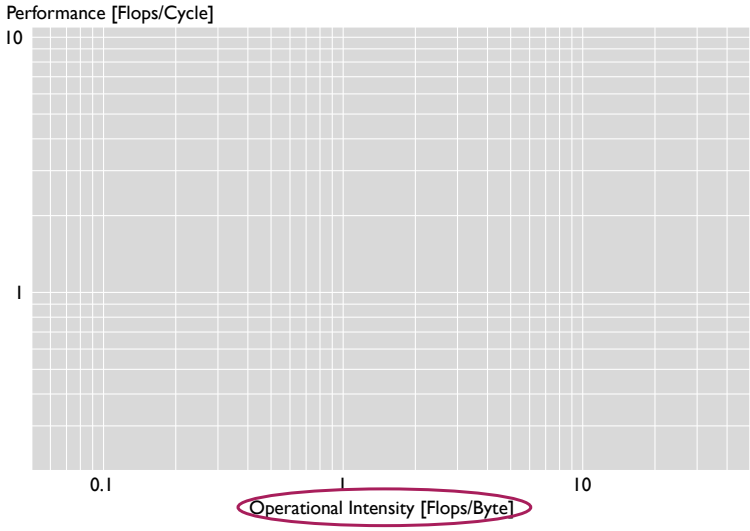

[Williams, 2009] "Roofline: An Insightful Visual Performance Model for Multicore", S. Williams *et al.* Communications of the ACM, 2009

Roofline Plot

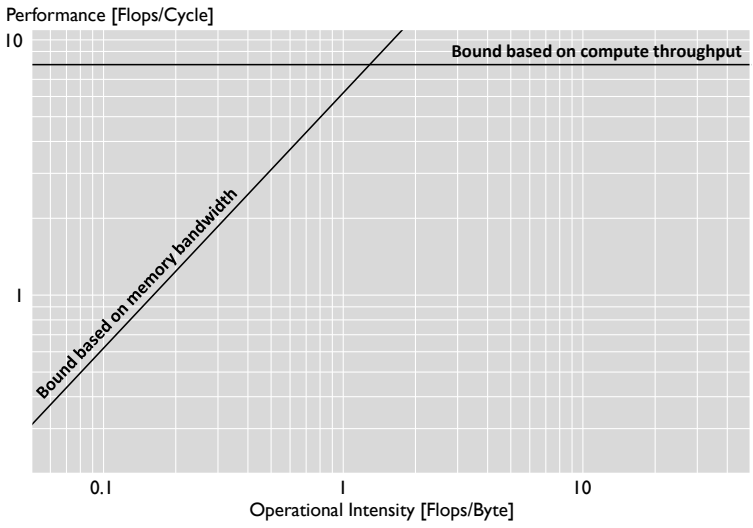

[Williams, 2009] "Roofline: An Insightful Visual Performance Model for Multicore", S. Williams *et al.* Communications of the ACM, 2009

Roofline Plot

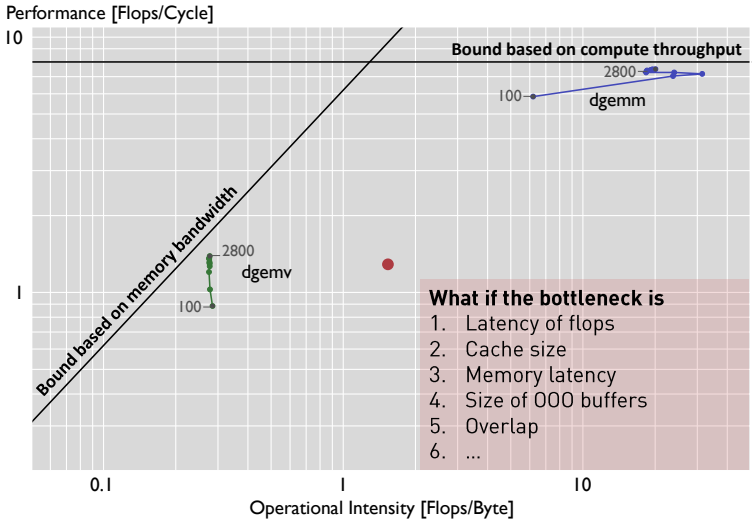

Bottleneck Modelling

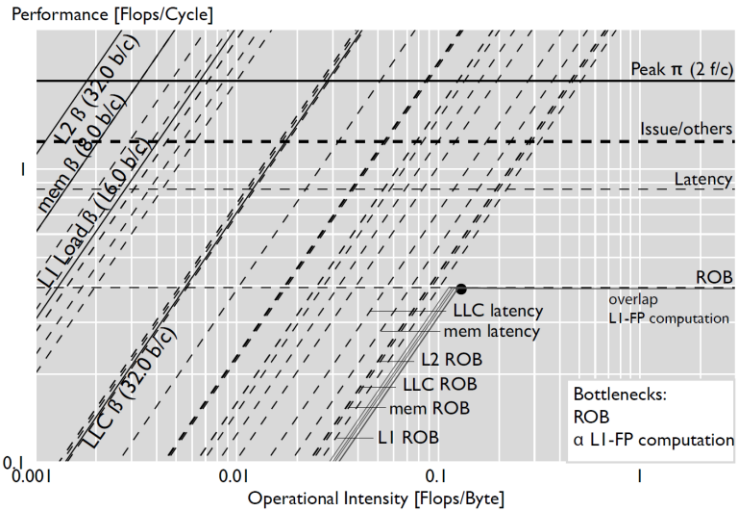

Example FFT 2^{20}

Adapting the Processor

FFT 1024

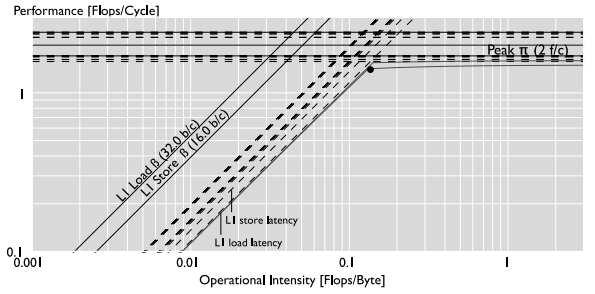

After adapting

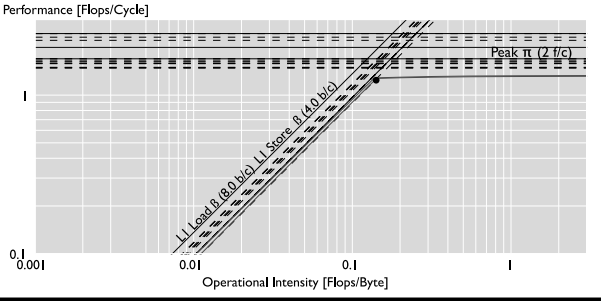

Evolution of Performance

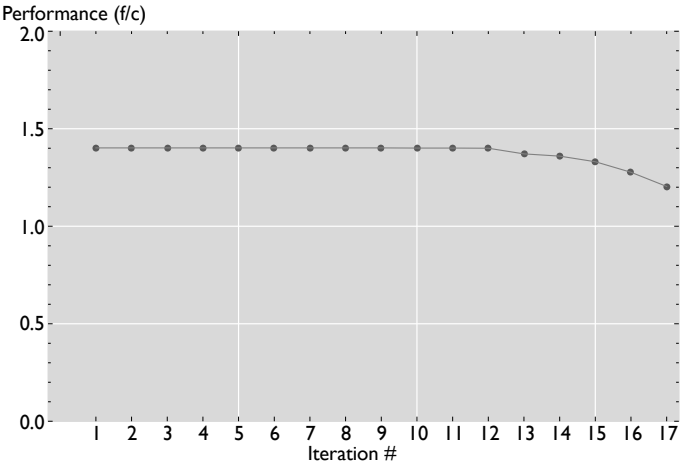

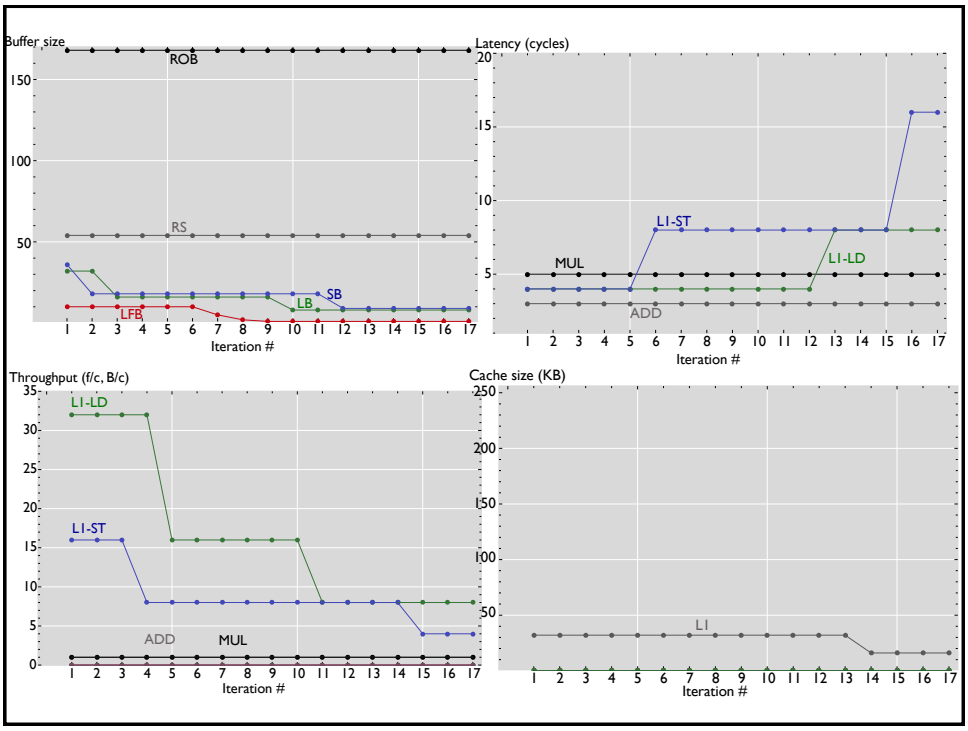

1. **Program synthesis for performance**
   Daniele Spampinato [CGO 2014]
2. **Performance bottleneck modeling**
   Victoria Caparròs Cabezas [IISWC 2014]
3. **Predicting Pareto-optimal solutions**
   Marcela Zuluaga, Andreas Krause [ICML 2013]

Multi-Objective Optimization

Case Study: Different Hardware Implementations of a 256-input Sorter

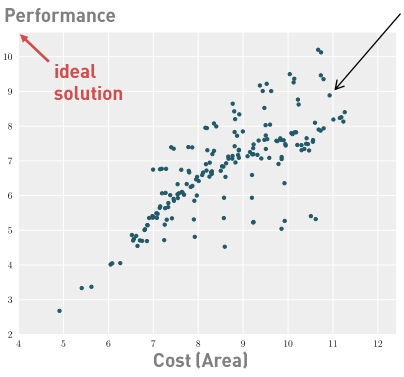

each determined by a set of 4 parameters

Which ones are relevant?

Multi-Objective Optimization

Case Study: Different Hardware Implementation of a 256-input Sorter

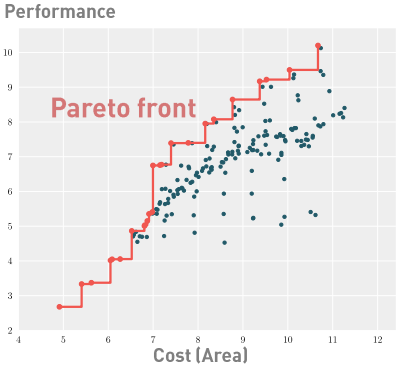

How to get it?

Multi-Objective Optimization

Case Study: Different Hardware Implementation of a 256-input Sorter

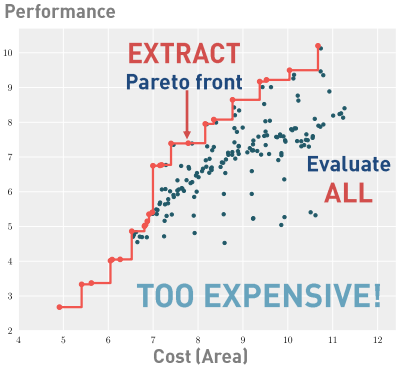

Multi-Objective Optimization

Case Study: Different Hardware Implementation of a 256-input Sorter

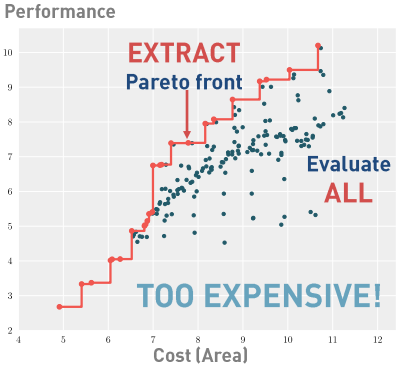

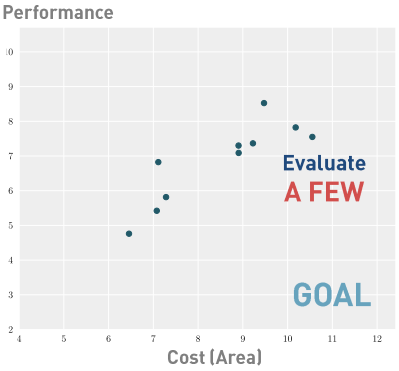

Multi-Objective Optimization

Case Study: Different Hardware Implementation of a 256-input Sorter

Our Solution: Pareto Active Learning (PAL)

Running PAL

Modeling with Gaussian Processes

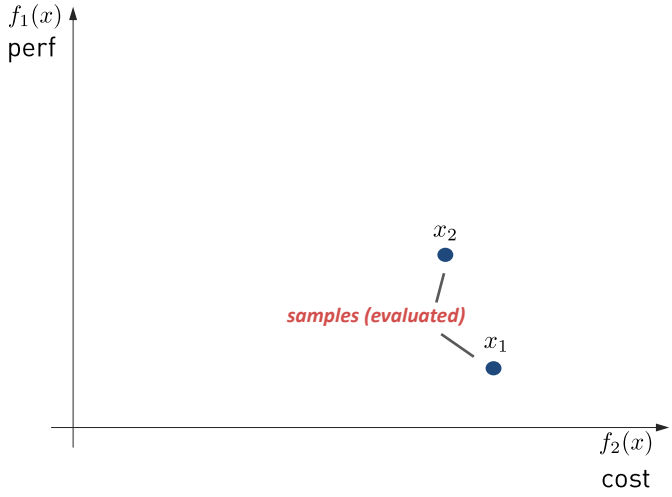

Running PAL

Modeling with Gaussian Processes

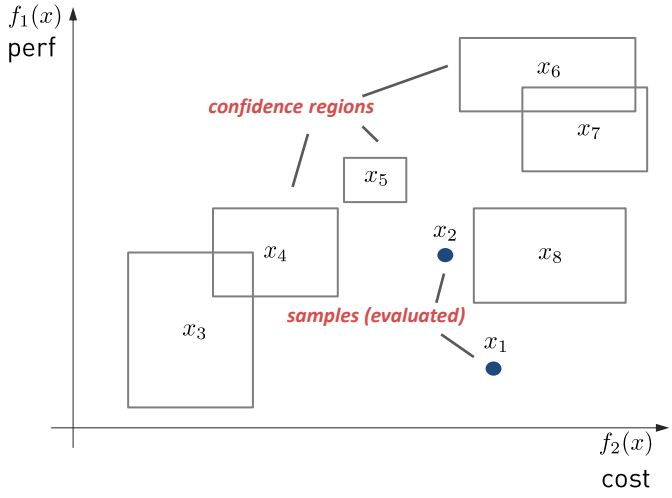

Running PAL

Classification

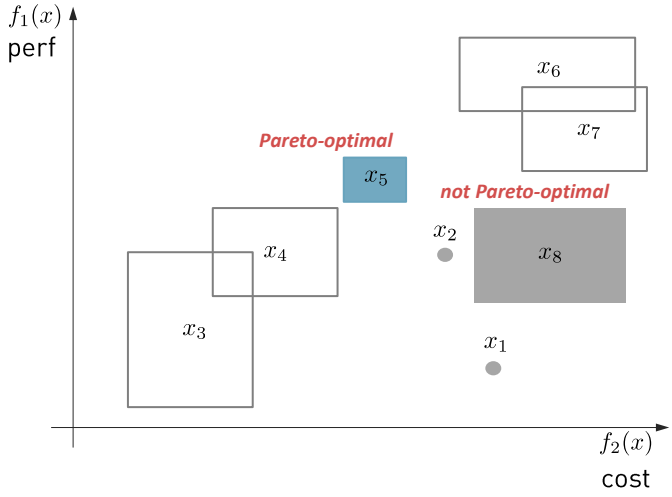

Running PAL

Sampling

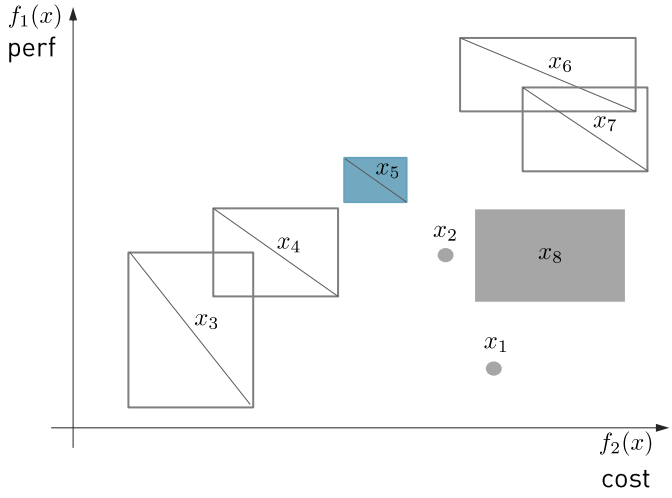

Running PAL

Sampling

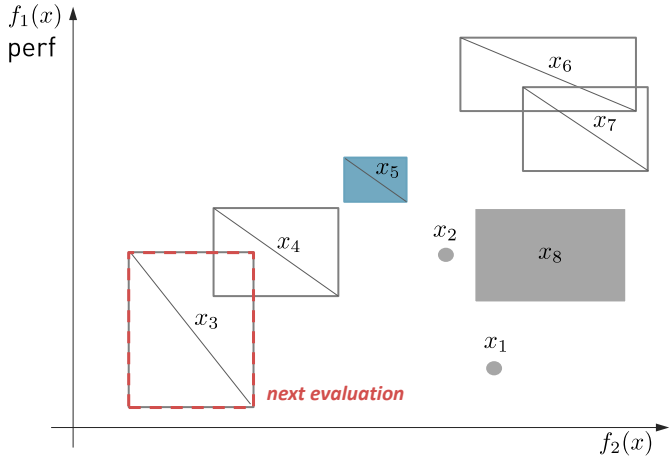

Modeling, Classification, Sampling

until all classified

Convergence Behaviour

ParEGO: State-of-the-art evolutionary algorithm

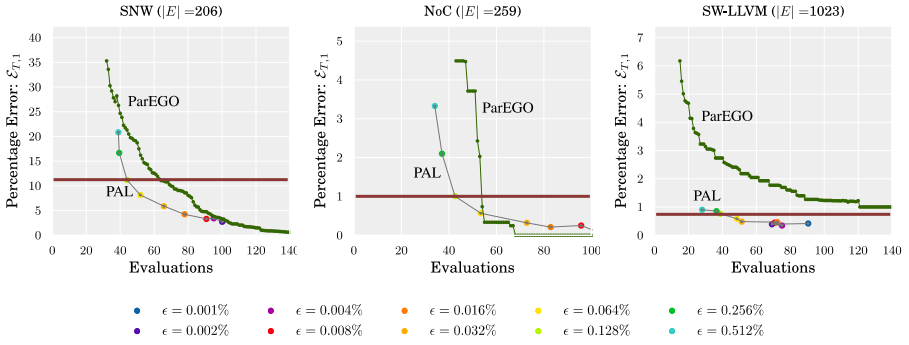

Program synthesis for performance

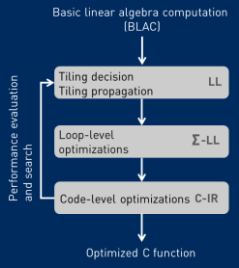

Bottleneck modelling

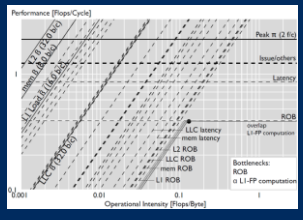

Predicting Pareto fronts

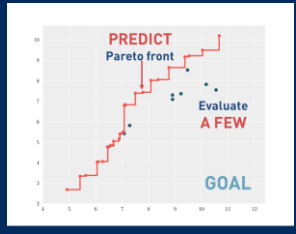

I am happy to discuss more
www.acl.inf.ethz.ch

Experiments: Data sets

4 features

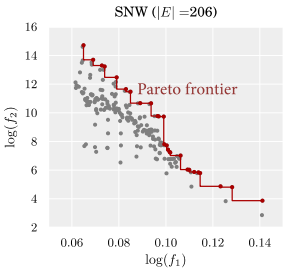

4 features

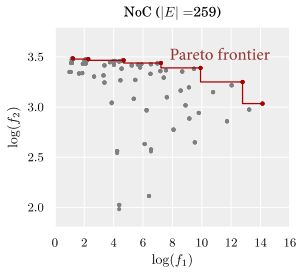

11 features

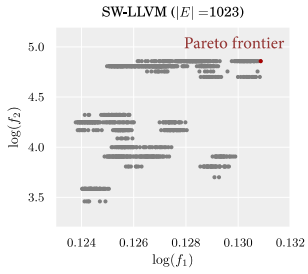

Marcela Zuluaga, Andreas Krause, Peter Milder, Markus Püschel. *Streaming Sorting Networks. DAC 2012*

Oscar Almer, Nigel Topham, Björn Franke. *A Learning- Based Approach to the Automated Design of MP-SoC Networks. ARCS 2011*

Predicting Performance via Automated Feature-Interaction Detection . *N. Siegmund, S. S. Kolesnikov, C. Kastner , S. Apel, D. Batory, M. Rosenmuller, and G. Saake. ICSI 2012*